

Mandatory Access Control in PostgreSQL - giving users ownership of their data

Leon du Toit

2019-01-15

Outline

- ▶ why take data ownership seriously?
- ▶ why Mandatory Access Control?
- ▶ a brief introduction to the `pg-need-to-know` module
- ▶ a use case to demonstrate features:
 - ▶ For users: ownership, insight and consent-based usage
 - ▶ For administrators: fine-grained access control, audit information
 - ▶ For developers: a rich REST API, with a built-in authorization model

Why take data ownership seriously?

- ▶ Regulations of the GDPR
 - ▶ increased focus on data privacy and protection
 - ▶ right to access
 - ▶ right to be forgotten
 - ▶ data portability
 - ▶ consent-based data usage
 - ▶ increased demand for audit information
- ▶ To counter surveillance capitalism
 - ▶ you (and your data) are the product
 - ▶ building applications to fight this trend

What Mandatory Access Control?

- ▶ *enforcible* policies, in contrast to Discretionary Access Control
- ▶ enables consent-based data access
- ▶ supports granular access needs

pg-need-to-know

- ▶ PostgreSQL “module” - really just a set of tables, views, and functions
- ▶ implements Mandatory Access Control
- ▶ more limited approach than SEPostgreSQL
- ▶ source: <https://github.com/leondutoit/pg-need-to-know>
- ▶ written in PL/pgSQL
 - ▶ procedural language, extending SQL with control structures
 - ▶ used to create functions
 - ▶ ~1000 sloc, another ~1500 for tests
- ▶ designed to be used via a REST API

Use case

Key terms:

- ▶ data owner: provides data about themselves
- ▶ data user: analyses data about others
- ▶ admin: creates and implements access control policies

Use case

Assume the following setup:

data owners: A, B, C, D, E, F

tables: t1, t2, containing data from all data owners

data users: X, Y, Z

Use case

Now suppose we need to set up the following access control rules in our DB:

- ▶ data users X, and Y should only have access to data in tables t1 and only data from owners A, B, C, D
- ▶ data user Z should have access to all data - i.e. tables t1, t2

Use case

Using pg-need-to-know, we implement this with the following groups, and table grants:

group1

- members: ((X, Y), (A, B, C, D))
- select table access grant: (t1)

group2

- members: ((Z), (A, B, C, D, E, F))
- select table access grants: (t1, t2)

pg-need-to-know via REST

HTTP client -> webapp -> REST server -> (pg-need-to-know, I