

Mandatory Access Control in PostgreSQL - giving users ownership of their data

Leon du Toit

2019-01-15

Outline

- ▶ why take data ownership seriously?
- ▶ why Mandatory Access Control?
- ▶ a brief introduction to the `pg-need-to-know` module
- ▶ a use case to demonstrate features:
 - ▶ For users: ownership, insight and consent-based usage
 - ▶ For administrators: fine-grained access control, audit information
 - ▶ For developers: a rich REST API, with a built-in authorization model

Why take data ownership seriously?

- ▶ Regulations of the GDPR
 - ▶ increased focus on data privacy and protection
 - ▶ right to access
 - ▶ right to be forgotten
 - ▶ data portability
 - ▶ consent-based data usage
 - ▶ increased demand for audit information
- ▶ To counter surveillance capitalism
 - ▶ you (and your data) are the product
 - ▶ building applications to fight this trend

What Mandatory Access Control?

- ▶ *enforcible* policies, in contrast to Discretionary Access Control
- ▶ enables consent-based data access
- ▶ supports granular access needs

pg-need-to-know

- ▶ PostgreSQL “module” - really just a set of tables, views, and functions
- ▶ implements Mandatory Access Control
- ▶ more limited approach than SEPostgreSQL
- ▶ source: <https://github.com/leondutoit/pg-need-to-know>
- ▶ written in PL/pgSQL
 - ▶ procedural language, extending SQL with control structures
 - ▶ used to create functions
 - ▶ ~1000 sloc, another ~1500 for tests
- ▶ uses Row-Level Security policies to implement MAC
- ▶ designed to be used via a REST API

Use case

Key terms:

- ▶ data owner: provides data about themselves
- ▶ data user: analyses data about others
- ▶ admin: creates and implements access control policies

Use case

Assume the following setup:

- ▶ data owners: A, B, C, D, E, F
- ▶ data users: X, Y, Z
- ▶ tables: `spending_habits`, `personal_details`, containing data from all data owners

Use case

Now suppose we need to set up the following access control rules in our DB:

- ▶ data users X, and Y should only have access to data in tables spending_habits and only data from owners A, B, C, D
- ▶ data user Z should have access to all data - i.e. tables spending_habits, personal_details

Use case

A hypothetical sequence of events:

1. admin creates tables
2. data owners and data users register themselves, data is collected
3. admin creates groups, adds members, adds table grants
4. data is analysed
5. users manage their own data
6. admins get audit insights
7. developers create applications using composing these features

Table creation

```
set role admin_user;
SET
select table_create(
  '{"table_name": "spending_habits",
    "columns": [
      {"name": "spending", "type": "int",
        "description": "Amount spent in NOK"},
      {"name": "item_type", "type": "text",
        "description": "Type of item purchased"},
      {"name": "purchase_date", "type": "date",
        "description": "Year-Month-Day on which purchase occurred"} ],
    "description": "data about spending habits"}'::json,
  'mac');
table_create
-----
Success
(1 row)
```

Figure 1:Creating a new table

User registration

- ▶ can require consent before user registration
- ▶ data collection not possible without registration

Group setup, table grants

- ▶ can link consent(s) to groups via group metadata
- ▶ group1
 - ▶ members: ((X, Y), (A, B, C, D))
 - ▶ select table access grant: (spending_habits)
- ▶ group2
 - ▶ members: ((Z), (A, B, C, D, E, F))
 - ▶ select table access grants: (spending_habits, personal_details)

Data analysis

```
set session "request.jwt.claim.user" = 'user_X';
SET
select current_setting('request.jwt.claim.user');
current_setting
-----
user_X
(1 row)

select * from spending_habits;
      row_id      | row_owner | row_originator | spending | item_type | purchase_date
-----+-----+-----+-----+-----+-----
4ad3b11e-32ff-42a1-850c-aff1f93f190e | owner_A  | owner_A        |      140 | food      | 2019-01-02
975f2758-5749-4915-bac1-48530f703062 | owner_A  | owner_A        |      100 | drink     | 2019-01-03
899efca4-a935-4e0d-ba25-29c4413d7c2a | owner_B  | owner_B        |       60 | drink     | 2019-01-02
7ef73351-1e7d-4f26-989b-d55fa0f1bfa5 | owner_B  | owner_B        |       78 | drink     | 2019-01-04
1be698b3-e1c0-4b98-a236-f273883f67dc | owner_C  | owner_C        |     1020 | travel    | 2019-01-04
c225db92-2171-4d21-9b7a-67c4ef0ad942 | owner_C  | owner_C        |      101 | food      | 2019-01-04
123f6322-130a-4b13-8f2a-2dbe0f0a9523 | owner_D  | owner_D        |      230 | travel    | 2019-01-05
ca97c462-7fea-49ce-8bde-fcef08a910ab | owner_D  | owner_D        |      448 | travel    | 2019-01-06
(8 rows)

select * from personal_details;
psql:./src/11-user-X-data-access.sql:7: ERROR:  access denied to table
CONTEXT:  PL/pgSQL function ntk.data_user_group_membership_with_correct_privileges(uuid,text,text) line 18 at
RAISE
```

Figure 2: User X's data access

Data analysis

User Z's data access

Data ownership

- ▶ right to access
- ▶ data portability
- ▶ right to be forgotten

Right to access

Owner A's data access

Data portability

- ▶ owner A can simple download their data

Right to be forgotten

Owner B deletes their data

Audit insights

- ▶ data access
- ▶ access control changes
- ▶ user initiated group removals
- ▶ user initiated data deletions
- ▶ data updates

Audit: data access

Audit: access control changes

Audit: user initiated group removals

Audit: user initiated data deletions

Audit: data updates

Application development

Architecture:

`client -> webapp -> REST -> (pg-need-to-know, PostgreSQL)`

- ▶ developers can focus on business logic
- ▶ authorization taken care of